

COMP3911 Secure Computing

18: Intrusion Detection & Incident Response


Nick Efford

<https://comp3911.info>

Today's Objectives

- For you to understand the basic principles of intrusion detection and see examples of how network-based intrusion detection can be done
- For you to appreciate the role of honeypots in detecting intrusion and studying attacker behaviour
- For you to learn two approaches to collecting evidence of attack from a computer system
- For you to see how we analyse a suspicious executable

The Challenge



Approaches



UNIVERSITY OF LEEDS

- Statistical anomaly detection
 - Develop profiles of each user's activity over time
 - Perform statistical tests to check whether deviation from observed profiles is significant
- Rule-based detection
 - Rules for deviation from previous usage patterns...
 - ... or for explicit suspicious behaviour

Network-Based IDS

These monitor network traffic for suspicious events...



<https://www.snort.org>



<https://zeek.org>

- Signature-based, cross-platform IDS
- Can sniff packets in real time or operate on previously captured traffic
- Preprocessors for IP defragmentation, TCP stream reassembly, portscan detection...
- Supports a very wide range of rules to identify attacks by their signature
 - New rules can appear within hours of an exploit becoming public knowledge!

Rule Examples

datagrams from any IP
address & any port...

... to subnet
192.168.5...

... on port used by
Back Orifice RA Trojan

```
alert udp any any -> 192.168.5.0/24 31337 \  
( msg: "Back Orifice"; )
```

```
alert tcp any any -> any any \  
( msg: "Shellcode?"; content "|90 90 90 90|"; )
```

What problem might there
be with this example?

TCP stream has four
consecutive values of 0x90
(NOP on Intel CPUs)

Hiding a NOP Sled



UNIVERSITY OF LEEDS

```
00000000: 9090      nop
00000002: 91        xchg      ecx,eax
00000003: 91        xchg      ecx,eax
00000004: 9090      nop
00000006: 91        xchg      ecx,eax
00000007: 91        xchg      ecx,eax
00000008: 9090      nop
0000000A: 91        xchg      ecx,eax
0000000B: 91        xchg      ecx,eax
0000000C: 9090      nop
0000000E: 91        xchg      ecx,eax
0000000F: 91        xchg      ecx,eax
00000010: 81E086FFFAF2 and      eax,0F2FAFF86 ;'≥· å'
00000016: B964010000 mov      ecx,000000164 ;' 0d'
0000001B: 29CC      sub      esp,ecx
0000001D: 33D2      xor      edx,edx
0000001F: 87E7      xchg      edi,esp
00000021: 89FC      mov      esp,edi
00000023: 81E0E1A3D9A3 1and     eax,0A3D9A3E1 ;'ú¹ úß'
```

Two successive register swaps have no net effect on behaviour, but break up the sequence of 0x90 values found in traditional NOP sleds...

Host-Based IDS

Example: [Tripwire](#)

- Creates and stores secure checksums of critical files
(How do you think it does this?)
- Comparison with these checksums allows unexpected changes to be detected
- Only detects impact of intrusion; can't detect attacker's recon probes, for example

Honeypots

- Security resources whose value lies in being probed, attacked or compromised
- Production honeypots
 - A specialised form of IDS
 - Used for incident response (evidence collection)
- Research honeypots
 - Capture of automated threats – e.g., worms
 - Observation of new tools / techniques
 - Investigation of attacker behaviour

Interaction Level

Low interaction

- Listen on standard ports
- May give a limited, superficially believable response
- Examples: Specter, Honeyd

High interaction

- Provide or simulate full-blown OS and applications
- Good for learning about attacker behaviour
- High risk: what if attacker breaks free?

Benefits & Drawbacks

Benefits

- Smaller quantities of higher-quality data, compared with firewall logs or NIDS alerts
- Simple to set up: no special rules, no signature database to maintain, etc

Drawbacks

- Narrow field of view
- Difficulty of accurate simulation
- Risk of honeypot being used to mount attacks

S P E C T E R
Engine Version : R 8.00
Threads : 17
Connections so far : 0

Operating System

- ☐ Random
- ☐ Windows 98
- ☐ Windows NT
- ☐ Windows 2003
- ☒ Windows XP
- ☐ MacOS
- ☐ MacOS X
- ☐ Linux
- ☐ Solaris
- ☐ NeXTStep
- ☐ Tru64
- ☐ Irix
- ☐ Unisys Unix
- ☐ AIX
- ☐ FreeBSD

Services

- ☒ FTP ?
- ☒ TELNET ?
- ☒ SMTP ?
- ☒ FINGER ?
- ☒ HTTP ?
- ☒ NETBUS ?
- ☒ POP3 ?
- ☒ Provide mails

Traps

- ☒ DNS ?
- ☒ IMAP4 ?
- ☒ SUN-RPC ?
- ☒ SSH ?
- ☒ SUB-7 ?
- ☒ BO2K ?
- ☒ GENERIC ?

Generic Trap Name: IRC

Generic Trap Port: 6667

Notification

- ☒ Incident DB ?
- ☒ Alert mail ?
- ☒ Short mail ?
- ☒ Status mail ?
- ☒ Event log ?
- ☐ Syslog ?

Configure Syslog

Intelligence

- ☒ Finger ?
- ☒ Trace Finger ?
- ☒ Port Scan ?
- ☒ DNS Lookup ?
- ☒ Whois ?
- ☒ Telnet Banner ?
- ☒ Ftp Banner ?
- ☒ SmtP Banner ?
- ☒ Http Header ?
- ☒ Http Document ?
- ☒ Trace Route ?

Max. Hops: 30

Character

- ☐ Random
- ☐ Failing ?
- ☐ Secure ?
- ☒ Open ?
- ☐ Aggressive ?
- ☐ Strange ?

Password Type

- ☐ Easy ?
- ☒ Normal ?
- ☐ Hard ?
- ☐ Mean ?
- ☐ Fun ?
- ☐ Cheswick ?
- ☐ Warning ?

☒ Send PW file ?

Watcher Setup

R

Vulnerability DB update installed (4897 bytes) [Fri Jan 27 22:19:34 2009]
Content DB is up-to-date [Fri Jan 27 22:19:46 2009]

FTP	running
TELNET	running
SMTP	running
FINGER	running
HTTP	running
NETBUS	running
DNS	running
SUB-7	running
SUN-RPC	running
POP3	running
IMAP4	running
BO2K	running
SSH	running
GENERIC	running

Engine Messages

☒ Errors ☒ Connections

Start Engine
Reconfigure
Load
About

Stop Engine
Log Analyzer
Save
License

Host Name : athena.mit.edu ?
System Name : OUTPOST ?
Configuration Version : 1.0 ?
Mail Server IP Address : 192.168.1.250 ?
Mail Address : admin@specter.com ?
Short Mail Address : inci@specter.com ?

User Configuration ?
Network Configuration ?
Web Service Configuration ?

☒ Include settings in mails ?
Status Mail Period [h] : 24 ?

☒ Remote Management Port : 28
Set Password ?

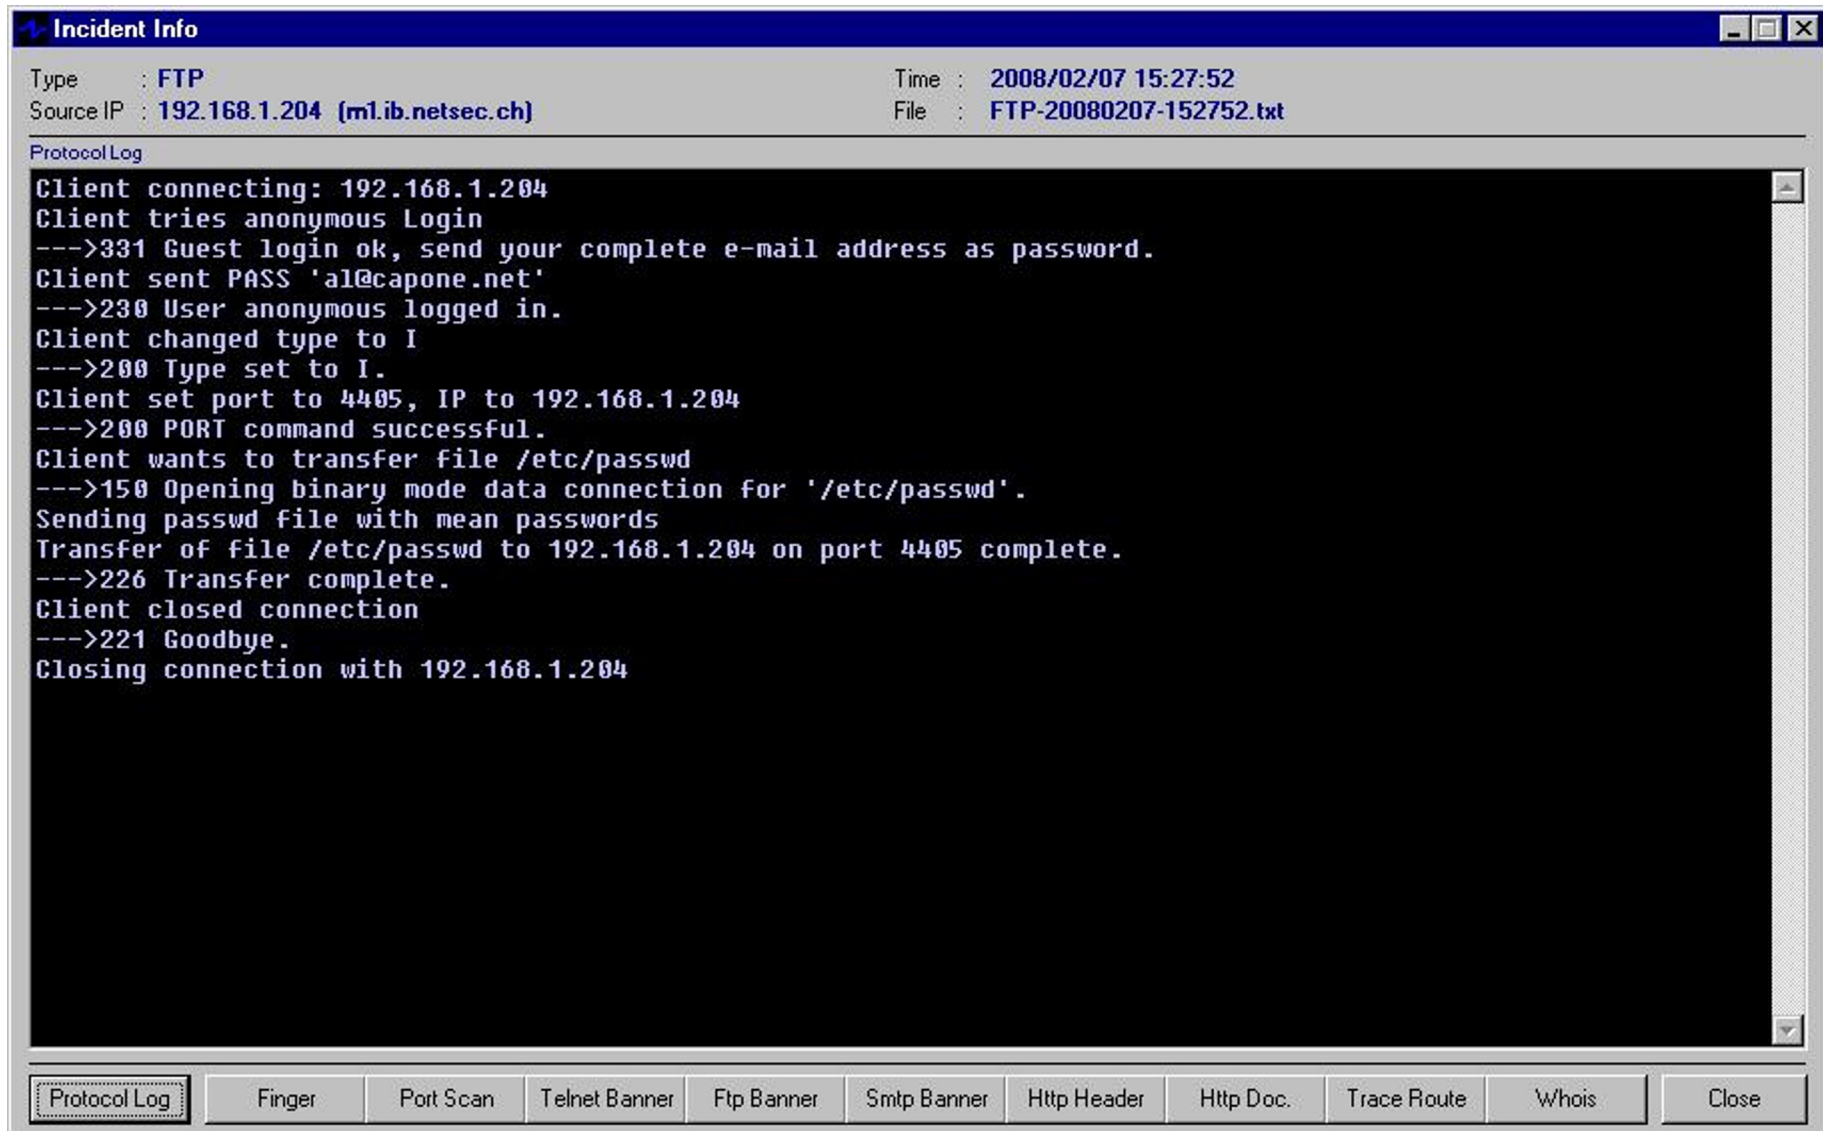
☒ Expect friendly connections
IP Addresses ?

☒ Use custom mail message for POP3
Edit Message ?

☒ Use custom warning message ?

Your actions are logged, intrusion alert was activated!

13



Incident Response

- Penetrated systems are subject to **forensic analysis**
- Goals are to
 - Discover how attackers gained access
 - Deduce what they were able to do on the system
 - Possibly gather evidence of criminal activity that is solid enough to use for prosecutions

Approaches



UNIVERSITY OF LEEDS

- Post-mortem analysis
 - Non-volatile data
- Live response
 - Volatile data
 - Non-volatile data
- Network-based evidence collection

Post-Mortem Analysis Steps

1. Forensic duplication
2. Recovery of deleted files
3. Collection of MACtimes and hashes
4. Removal of known files
5. Identification of files by signature
6. Reverse engineering of unknown executables
7. Reconstruction of activity (email, browsing, etc)

Forensic Duplication

Common to remove evidence drives in order to duplicate them

Why do we use dedicated connectors when duplicating?
Why not plug the evidence drives directly into the forensic PC?

Connectors allow hot-swapping and, more importantly, provide **hardware-based write blocking**



Recovery of 'Deleted' Data

- For most OS, deletion merely breaks link between file metadata and disk blocks holding file contents
 - ... and marks those blocks as unused
- So **data persists on disk** until those blocks are reused!
 - Persistence times depend upon number of free blocks available and level of system activity
 - Farmer & Venema found that 'data half-life' varied between 12 and 35 days for small servers

File MACtimes

- UNIX-like filesystems define three file timestamps:
 - Last modification time (`mtime`)
 - Last access time (`atime`)
 - Time of last status change (`ctime`)
- Equivalents exist in other OS (Windows NTFS, etc)
- MACtimes provide useful evidence of system activity...
- ... which can be easily lost by reboots, system backup, etc!

Computation of File Hashes

- Do it for
 - Entire duplicated disk partition
 - Individual files recovered from that partition
- UNIX-based systems offer multiple command line tools that can be used for this, e.g.:

```
shasum -a 256 foo.txt
```

- But why do we need this?

Removal of Known Files

- Typical disk may have on the order of 10^5 files, only a few of which will be interesting to us
- Hash databases for OS programs, standard system files and common apps can be used to ‘narrow the field’
 - Example: [NIST National Software Reference Library](#)

Static Analysis of Executables

- Extract text strings
- Perform a hex dump
- View symbol information
- View shared objects (DLLs)
- Examine binary file format
- Disassemble code

Dynamic Analysis of Executables

- Use VMware or similar virtualisation software to run the unknown executable in a secure sandbox
- Trace system/library calls
- Step through the code in a debugger

Live Response: Volatile Data

- System date & time
- Running processes, process memory dumps
- Running services (Windows)
- Loaded kernel modules (Linux)
- Scheduled jobs
- Network connections & open ports
- Executables that have opened ports
- Opened files
- Users currently logged on

Live Response: Approach

- Must adhere to the **Order Of Volatility Principle** – capture most volatile data first, least volatile last
- Evidence collection tools must be trustworthy and self-contained
- Typical approach is to run the tools directly from a CD containing executables & libraries
- Data can be transferred from victim to forensic workstation via the network, using `netcat`
 - Hash of any captured evidence should be computed immediately

Summary

We have

- Considered some of the principles behind automated intrusion detection
- Explored the use of network-based and host-based IDS
- Discussed how to conduct a post-mortem analysis of an attacked computer system
- Investigated an example of a suspicious executable

Follow-Up / Further Reading

- Chapters 5 & 16 of [*Thinking Security*](#)
- Provos, *Virtual Honeypots*, Addison-Wesley 2007
- Farmer & Venema, *Forensic Discovery*, Addison-Wesley, 2005